

```

In [4]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

df = pd.read_csv('Housing.csv')

print(df.head())

print(df.describe())

print(df.isnull().sum())

correlation_matrix = df.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title("Correlation Matrix")
plt.show()

X = df[['bedrooms', 'bathrooms',]]
y = df['price']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size

```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom
0	13300000	7420	4	2	3	yes	no
1	12250000	8960	4	4	4	yes	no
2	12250000	9960	3	2	2	yes	no
3	12215000	7500	4	2	2	yes	no
4	11410000	7420	4	1	2	yes	yes

	hotwaterheating	airconditioning	parking	prefarea	furnishing	stat
0	no	yes	2	yes	furnish	
1	no	yes	3	no	furnish	
2	no	no	2	yes	semi-furnish	

```

3          no          yes          3          yes          furnish
ed
4          no          yes          2          no          furnish
ed
          price          area          bedrooms          bathrooms          stor
ies \
count  5.450000e+02          545.000000          545.000000          545.000000          545.000
000
mean    4.766729e+06          5150.541284          2.965138          1.286239          1.805
505
std     1.870440e+06          2170.141023          0.738064          0.502470          0.867
492
min     1.750000e+06          1650.000000          1.000000          1.000000          1.000
000
25%     3.430000e+06          3600.000000          2.000000          1.000000          1.000
000
50%     4.340000e+06          4600.000000          3.000000          1.000000          2.000
000
75%     5.740000e+06          6360.000000          3.000000          2.000000          2.000
000
max     1.330000e+07          16200.000000          6.000000          4.000000          4.000
000

```

```

          parking
count  545.000000
mean    0.693578
std     0.861586
min     0.000000
25%     0.000000
50%     0.000000
75%     1.000000
max     3.000000

```

```

price          0
area           0
bedrooms       0
bathrooms      0
stories        0
mainroad       0
guestroom      0
basement       0
hotwaterheating 0
airconditioning 0
parking        0
prefarea       0
furnishingstatus 0
dtype: int64

```

```


/var/folders/l9/49s71scs52x36hwkp6gmhklr0000gn/T/ipykernel_3135/13
23379855.py:22: FutureWarning: The default value of numeric_only i
n DataFrame.corr is deprecated. In a future version, it will defau
lt to False. Select only valid columns or specify the value of num
eric_only to silence this warning.

```

```

correlation_matrix = df.corr()

```



In [5]:

```
model = LinearRegression()

model.fit(X_train, y_train)
```

Out [5]:

```
▼ LinearRegression
LinearRegression()
```

In [6]:

```
# Model Evaluation
y_pred = model.predict(X_test)

mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("Mean Squared Error:", mse)
print("R-squared:", r2)
```

```
Mean Squared Error: 3737500333041.07
R-squared: 0.2605701076698277
```

In []:


In [8]: *# Predictions and Visualization*

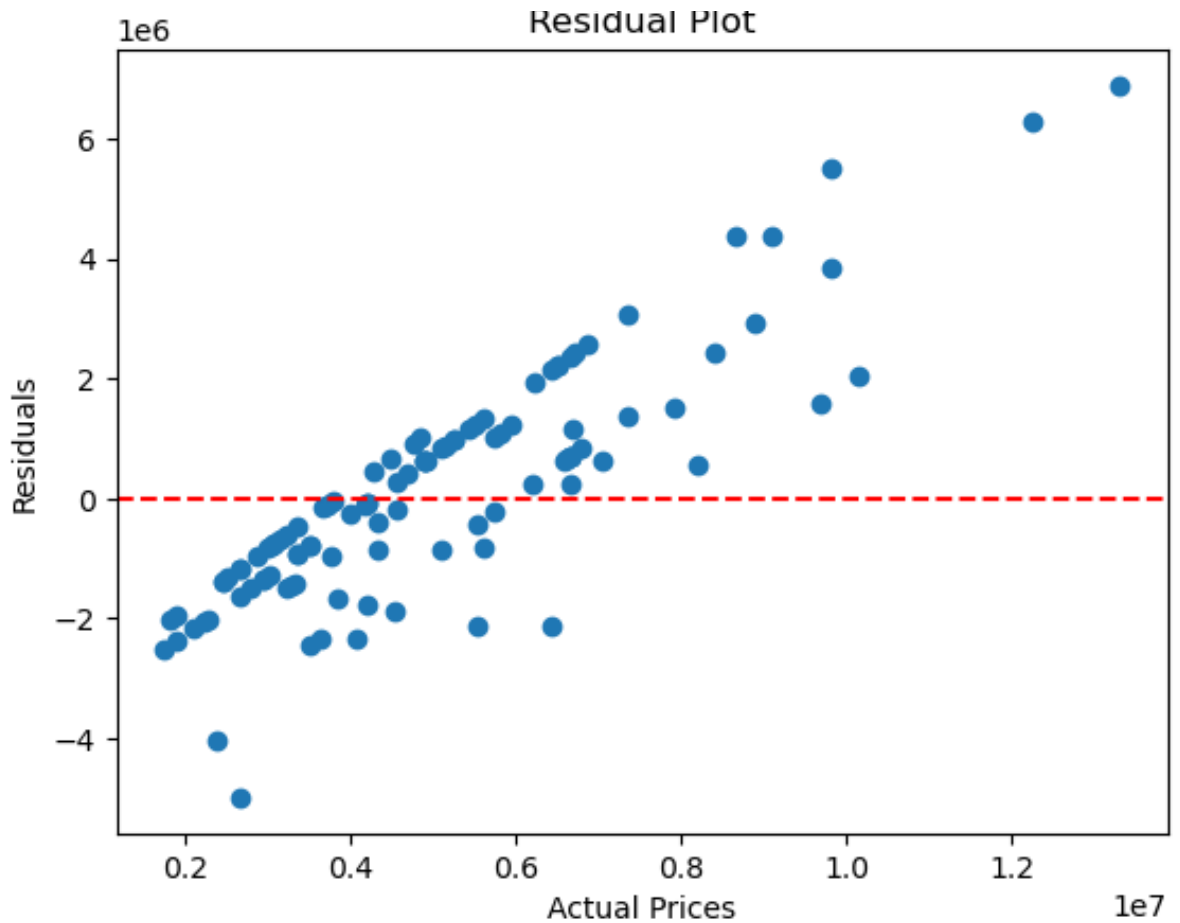
```
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual Prices vs. Predicted Prices")
plt.show()
```

```
residuals = y_test - y_pred
plt.scatter(y_test, residuals)
plt.axhline(y=0, color='red', linestyle='--')
plt.xlabel("Actual Prices")
plt.ylabel("Residuals")
plt.title("Residual Plot")
plt.show()
```

```
new_data = [[4000,3]]
predicted_price = model.predict(new_data)

print("Predicted Price:", predicted_price[0])
```





Predicted Price: 1816016323.403072

```
/Users/amjaf/anaconda3/lib/python3.11/site-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
  warnings.warn(
```

In []: